

Christopher Wood
Reflection Essay #5

The *expression problem* describes the programming language design issue of being able to define a program (or subset of the program) as a datatype that has a set of cases and associated processes (or functions) that operate on the datatype [2]. The datatype cases and supported processes should both be extensible while maintaining the strong static type safety of the datatype (i.e. processors are only allowed to use the datatype in the way described by their corresponding datatype variant/case). Furthermore, this extension should be supported after compilation in such a way that it does not affect any existing compiled code. As stated by Wadler, the ability for a language to solve the expression problem is a “salient indicator of its capacity for expression” [1].

Based on this description, we can see that this type of problem arises in situations where users may need to extend the operations applicable to built-in and user-defined datatypes. One example where this might occur is in the design and implementation of a graphics editing software tool. Although I am not speaking from experience, I would imagine there are cases where a user may want to define new shapes and transformation functions for these shapes in order to create different images. Thus, if we represent shapes (or more specifically, polygons) as a datatype with various cases in which it may be used, and then enable the user to specify new processors (i.e. transformation functions) that operate on this datatype, the user is free to “express” graphical images in novel ways. As an example, the user may want to define a processor for the polygon that treats it as a circle and transforms it into a semi-circle. We might also have a semi-circle polygon variant to which a user can define a duplicate processor, which in turn recreates a full circle from this semi-circle. There are undoubtedly many more examples where this problem comes into play in practical software development projects.

The utility of studying this problem depends on the perspective in which we view programming languages. If we consider a programming language to be strictly a form of expressing computations, then I would argue that the ability for a language to solve the expression problem becomes a very important indicator of its ability to fulfill this goal. Programmers tend to think about problems in very abstract and tangible ways, and so it would be ideal to not constrain them to built-in types defined at compile time. We naturally attempt to solve problems by reusing solutions to smaller problems. Thus, software solutions may be (and usually are) composed of expressions that solve sub-problems (developers do not like to “reinvent the wheel”). Granted, the developer could always revisit the software implementation and add new datatype variants or supported operations as needed. However, if this can be achieved at runtime with minimal computational overhead for resolving types and ensuring type safety, then we essentially lift the developer’s constraint to pre-defined datatypes and operations to solve specific problems. Consequently, the language promotes more natural problem solving through free-form expression.

Finally, it is interesting to note the *expression problem* solution differences in both the imperative and declarative programming language domains. Wadler’s solution is beautifully crafted in the imperative object-oriented paradigm. His application of the Visitor pattern, combined with generics, made it easy to add new cases to the datatype, but the addition of operations is somewhat convoluted. This seems to contradict the solutions presented in the

declarative functional domain (such as those implemented in Scala), which tend to provide easy ways to extend the processors for a datatype but make it difficult to include new cases. This key difference seems to provide further support for the fundamental differences between these two paradigms. Although both techniques present viable solutions to the problem, they both have their own implementation benefits, as is the usual case between imperative and declarative languages.

While my primary field of interest is not programming language design, I understand the importance and merit of “programming-languages problems” similar to the *expression problem*. Software engineers are often trained to solve problems within the context of a set of tools, but with novel solutions to these fundamental language design problems we can train future generations of engineers to think in terms of the problem at hand, and then simply find the best tool that fits the job.

References

- [1] P. Wadler. The expression problem. *java-genericity Mailing List*, 1998.
- [2] M. Zenger and M. Odersky. Independently extensible solutions to the expression problem. URL <http://scala.epfl.ch/docu/related.html>, 2004.